

What Makes AI a Good Cultural Mediator?

Evidence from Literary Paratexts

Lidia Zhou Mengyuan

Abstract

Generative AI systems increasingly introduce, explain, recommend, and reframe cultural works for new audiences, yet their evaluation is often framed as harm detection or generic text-generation quality. We propose an evaluation framework for AI as a cultural mediator, defining four operational dimensions of good mediation: availability, interpretive substance, cultural grounding, and discursive alignment. We apply the framework to AI-generated literary paratexts using 282 contemporary literary titles, 28,200 outputs, and five model families, with original-language and English generation as diagnostic conditions. The framework reveals a trade-off: English prompting reduces refusals from 64 to 24 and generic failures from 43 to 19, but also reduces average entity recall from 0.1819 to 0.1135 and average Jaccard overlap from 0.1169 to 0.0734. The result shows that gains in surface usability can coincide with losses in cultural grounding. For machine learning evaluation, the contribution is a task-grounded framework for diagnosing when dimensions of cultural mediation move in different directions.

1. Introduction

Generative AI is increasingly asked not only to produce culture, but also to mediate it. Systems now write book descriptions, explain artworks, summarize films, generate tourist guides, prepare educational materials, translate cultural references, and recommend cultural items to new audiences. In these settings, the model does not simply output text. It selects what matters, omits some details, and frames a cultural object for users who may not share its background.

This mediating role calls for a richer account of evaluation. Much existing work on cultural aspects of AI focuses on avoiding failure: reducing harmful stereotypes, misinformation, bias, and value misalignment (Adilazuarda et al., 2024; Alkhamissi et al., 2024; Pawar et al., 2025). Those concerns remain essential, but they do not define what successful cultural AI looks like. A system may avoid obvious harm and still offer poor cultural mediation: it may be fluent but empty, accessible but flattening, or aligned with familiar discourse while losing the local details that make a cultural object situated.

We therefore ask: *what makes AI a good cultural mediator, and how can this be evaluated in a way that is useful for machine learning research?* We propose a four-part framework: availability, interpretive substance, cultural grounding, and discursive alignment. Each dimension is linked to an observable evaluation procedure: refusal detection, generic-failure annotation, multilingual entity retention, and semantic similarity to reference discourse. These dimensions are not isolated components. They may interact dynamically: improving the threshold ability to mediate may weaken the contextual specificity of mediation. The goal is not to collapse cultural mediation into a single number, but to decompose it into dimensions that make such tensions measurable.

We demonstrate the framework through AI-generated literary paratexts: short public-facing texts such as publisher descriptions, jacket-copy style introductions, and platform summaries. Literary paratexts are a useful case because they are bounded, culturally meaningful, and comparable across models. They sit between factual description and interpretation. A good paratext should be readable, but also specific; accessible, but not flattening; cautious, but not empty. This makes the task a tractable controlled evaluation setting for cultural AI evaluation, while still preserving a genuinely interpretive function. This motivation aligns with recent calls for evaluating generative AI as a cultural and interpretive technology rather than only as a text generator (Kommers et al., 2026; Qadri et al., 2025).

Our case study uses 282 contemporary literary works from 11 regional groupings and 28,200 publisher-style introductions generated by five model families: Claude Haiku 4.5,

Department of Translation, The Chinese University of Hong Kong. Correspondence to: Lidia Zhou Mengyuan <lidi Zhou@cuhk.edu.hk>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

Gemini 3.1 Flash-Lite Preview, GPT-5 mini, Mistral Large, and DeepSeek Chat. We apply the framework under two diagnostic conditions: generation in the work’s original language and generation in English. This is not a prompt-optimization study. Rather, the contrast tests whether mediation closer to the work’s linguistic setting differs from mediation through English as a wider global language.

The central empirical finding functions as a validation case for the framework. English prompting improves availability and interpretive substance: refusals and generic failures decrease. But it weakens cultural grounding: entity retention declines across all model families. Without the framework, this would look like a simple improvement in usability. With the framework, it becomes visible as an English-prompting trade-off between mediating access and preserving cultural anchors.

The paper makes three contributions. First, it proposes a task-grounded framework for evaluating AI as a cultural mediator. Second, it operationalizes the framework using reproducible signals. Third, it demonstrates the framework’s diagnostic value through a contrast that reveals a trade-off hidden by single-metric evaluation. More broadly, it argues that cultural AI should be evaluated as a multi-objective problem in which mediation quality depends on relations among dimensions, not only on higher scores along one axis.

2. AI as Cultural Mediator

AI systems increasingly occupy positions within cultural circulation. They help users decide what a cultural object is, why it matters, and how it should be approached. These roles have long been associated with editors, teachers, translators, critics, curators, librarians, reviewers, and recommendation platforms. AI does not simply replace those actors, but its outputs now enter the chain of mediation through which cultural objects become visible, understandable, and valuable.

Cultural mediation involves more than information transfer. It includes selection, emphasis, contextualization, simplification, comparison, and orientation. A museum label does not merely list facts; it tells a visitor what to notice. A tourist guide frames a place through history, atmosphere, and identity. A book recommendation positions a title in relation to a reader’s interests and prior knowledge. In each case, mediation shapes reception. In this sense, AI-generated cultural descriptions can be understood as interpretive outputs: they organize context, foreground some meanings, and make others less visible (Kommers et al., 2026).

This is why cultural AI evaluation should not stop at harm mitigation. A system can avoid bias and misinformation and still perform cultural mediation poorly. It can make

an unfamiliar work appear accessible by translating it into generic language. It can remove names, places, and institutions because they seem less relevant to a broad audience. The output may be fluent and even useful, but culturally thin. A positive vision of cultural AI therefore requires an evaluation language for constructive success, not only for failure avoidance.

3. Framework

For model evaluation, the framework functions as an evaluation decomposition. Rather than asking whether a generated cultural description is simply “good,” it separates success into dimensions that can be measured, inspected, and compared across models, prompting conditions, and domains. Table 1 summarizes the operational signal and main diagnostic failure mode for each dimension. The four dimensions draw on a view of cultural mediation as a social and interpretive process: a mediator must first be able to intervene, must offer a meaningful orientation, must preserve links to cultural context, and must remain legible within a recognizable public genre. This matters because cultural mediation is a multi-objective problem: a system may improve on one objective while declining on another.

Formally, for a cultural object with metadata x and a target mediation condition c (e.g., language, audience, or platform), the model produces an introduction $y = f(x, c)$. The evaluation problem is not to assign y a single quality score, but to inspect a vector of criteria $m(y) = (m_1, m_2, m_3, m_4)$ corresponding to the four dimensions below. The framework is therefore diagnostic: it is designed to capture and quantify tensions among dimensions rather than treating them as independent checklist items.

Availability. Availability asks whether the model performs the requested mediating task at all. In literary paratext generation, an unavailable model refuses, gives only a disclaimer, or avoids producing an introduction. Conceptually, availability is the threshold condition for cultural intervention: without it, no mediation occurs. Technically, it captures abstention and refusal behavior that may otherwise be hidden when evaluation focuses only on successful generations.

Interpretive substance. Interpretive substance asks whether the output provides meaningful orientation. A good mediator should help an audience understand something about a work’s themes, setting, historical background, or significance. In cultural terms, this dimension corresponds to the intellectual responsibility of mediation: the output should not merely sound plausible, but should give the audience a specific reason for how to approach the work. Generic praise (e.g., “a powerful exploration of the human

Table 1. Operational view of the evaluation framework.

Dimension	Signal	Failure mode
Availability	Refusal count	Refusal/evasion
Substance	Generic-failure count	Empty output
Grounding	Recall; Jaccard	Context loss
Alignment	Semantic similarity	Discourse drift

condition”) is a common failure mode.

Cultural grounding. Cultural grounding asks whether the output preserves links between a cultural object and the specific worlds that give it meaning: people, places, communities, institutions, events, or histories. This dimension is the framework’s main guard against global genericity: the tendency to make cultural objects widely legible by detaching them from local anchors. Here we operationalize one aspect of grounding through multilingual named-entity retention. Entity-based grounding is only a proxy, but recent work on cultural grounding also treats culturally significant entities as a key site where models may misrecognize or lose context (Nyandwi et al., 2025). This measure does not capture everything—such as tone, form, or style—but it provides a scalable signal of whether concrete cultural anchors are retained.

Discursive alignment. Discursive alignment asks whether the output remains recognizable within public-facing forms of cultural introduction. A book description should resemble a plausible publisher-style introduction, while not necessarily copying existing descriptions. We operationalize this through semantic similarity to collected reference descriptions.

4. Case Study: Literary Paratexts

Literary works rarely circulate alone. They reach readers through titles, covers, blurbs, reviews, catalogues, metadata, publisher descriptions, and platform summaries. These materials often form the first point of contact between a reader and a book. Following Genette (1997), we treat such paratexts as mediating forms: they stand at the border of the work, but they shape how it enters circulation.

Publisher-style introductions are especially useful for applying the framework. They require the model to produce a short, coherent, accessible, and interpretive text. They are not as open-ended as fiction generation, where success is hard to define, but they are also not purely factual summaries. Good paratext is therefore a useful test bed for cultural AI evaluation.

The broader relevance extends beyond literature. Similar mediating functions appear when AI generates museum

labels, tourist guides, educational explainers, streaming recommendations, cultural heritage descriptions, translation notes, or summaries of international events. In each case, the system must decide what to foreground, what to omit, and how to make a cultural object legible to a new audience. Literary paratexts provide one controlled setting in which these broader issues become measurable.

5. Data and Methods

Corpus. The corpus contains 286 rows corresponding to 282 unique literary titles after title-level deduplication. The works were published between 2000 and 2025 and were selected from prize-winning or canonically recognized contemporary literature. The dataset spans 11 regional groupings: Argentina, Brazil, China, Egypt, Iceland, Nepal, Nordic countries, Serbia, Sub-Saharan Africa, the United Kingdom/English-language, and the United States. For each title, the dataset records bibliographic metadata and, where available, a short collected description used as a reference.

Task formulation. The task is controlled generation of a 120–180 word publisher-style introduction for a given literary work. Each prompt provides title, author, country or region, and original language, and instructs the model to remain general rather than invent plot details when uncertain.

Model families and diagnostic conditions. We compare five model families: Claude Haiku 4.5, Gemini 3.1 Flash-Lite Preview, GPT-5 mini, Mistral Large, and DeepSeek Chat. For each work, each model is run under two diagnostic conditions: generation in the work’s original language and generation in English. The contrast serves as a stress test for the framework. If the dimensions move together, the framework would mostly confirm a single overall quality pattern. If they diverge, the framework can reveal a trade-off that simpler evaluation would miss. Each title is generated 10 times per condition, producing 20 generations per title per model and 28,200 outputs overall.

Operationalization. Availability is measured through refusal counts. Interpretive substance is measured through generic-failure counts. Cultural grounding is measured through entity retention, using a spaCy-based multilingual NER pipeline applied consistently to reference and generated texts. We extract PERSON, GPE, LOC, NORP, EVENT, and ORG entities, deduplicate entity surface forms within each text, and compute overlap between reference and generated entity sets. For reference entities E_r and generated entities E_g , we report entity recall $|E_r \cap E_g|/|E_r|$ and Jaccard overlap $|E_r \cap E_g|/|E_r \cup E_g|$. Discursive alignment is measured through semantic similarity to collected reference descriptions. Our use of similarity is motivated by recent work cautioning against relying on a single text-generation

metric (Bhandari et al., 2020; Fabbri et al., 2021; Zhang et al., 2020; Yuan et al., 2021).

6. Framework Validation: Diagnostic Contrast

Availability and substance. Across all 28,200 outputs, we identify 88 refusals. All refusals are produced by Claude, suggesting a model-specific abstention pattern in this task. This is precisely what the availability dimension is designed to capture: a form of conservative or over-cautious model behavior that prevents mediation from taking place. Under the original-language condition there are 64 refusals out of 14,100 outputs (0.45%); under English there are 24 (0.17%). Generic failures show a similar pattern: 62 in total, with 43 under original-language generation and 19 under English. By the framework’s first two dimensions, English generation improves the model’s capacity to mediate: it is more likely to provide a usable, non-empty introduction.

Cultural grounding. The framework changes the interpretation once cultural grounding is considered. Across all five model families, original-language generation produces stronger entity retention than English generation. As shown in Figure 1, all five model families retain fewer cultural entities under English generation than under original-language generation, for both recall and Jaccard overlap. Averaged across models, entity recall falls from 0.1819 to 0.1135, while Jaccard overlap declines from 0.1169 to 0.0734. Comparing the availability and grounding dimensions therefore reveals an English-prompting trade-off that would not be visible from task completion alone: the system becomes a more available mediator while becoming a less culturally grounded one.

Discursive alignment. Semantic similarity to reference descriptions provides another useful signal. In the original-language condition, similarity is highest for Gemini at 0.6305, followed by DeepSeek at 0.6091, Mistral at 0.6034, GPT-5 at 0.6023, and Claude at 0.5823. However, alignment does not by itself settle whether mediation is good: an output may resemble a collected description while still losing culturally anchoring details.

What the framework reveals. The key finding is not that English generation is simply worse, nor that original-language generation is always better. Rather, the framework demonstrates its diagnostic value by showing that the dimensions of good mediation can move in different directions. A system can look better under one evaluative lens and worse under another. The English-prompting trade-off is therefore not an extra result added after the framework; it is the framework’s validation case. It shows how the evaluation decomposition can quantify a hidden trade-off between broader accessibility and weaker cultural grounding.




Figure 1. Aggregate diagnostic contrast. English generation improves availability and substance while reducing cultural grounding.

7. Discussion

The case study suggests that cultural AI should be evaluated as a multi-objective mediation problem. For model developers and evaluators, the framework provides a way to specify *which* dimensions of cultural performance are being optimized, and which may be weakened as a result. In our case study, English generation improves two dimensions of mediation while weakening a third. That kind of divergence is useful both for evaluation and for system design because it identifies when models become easier to deploy while becoming less capable of preserving cultural context.

The framework is intended to generalize beyond literary paratexts. In museum interpretation and cultural heritage, availability would ask whether the system can generate a useful label; substance would ask whether the label explains meaning rather than listing facts; grounding would include provenance, community affiliation, ritual use, or contested histories; alignment would concern fit with public museum discourse. In tourism and cross-cultural communication, grounding would include local names, historical layers, languages, and social practices. In education, it would include historical and linguistic context. In streaming recommendation, it would include genre traditions, regional contexts, or artistic lineages. In each domain, the operationalization would vary, but the evaluative logic remains stable: good AI mediation must make culture accessible without erasing the anchors that make it culturally specific.

The framework also supports a positive vision of cultural AI. Good mediation is not merely the absence of bias or misinformation; it is the presence of *situated access*: helping audiences approach cultural objects while preserving enough specificity for those objects to remain meaningful. This matters for cultural heritage preservation and cross-cultural circulation because AI-generated mediation may become the first layer through which users encounter a work, place, artefact, or tradition. At the same time, building AI systems that mediate culture more effectively is not an unambiguous good. More capable systems may reshape the work of editors, translators, teachers, curators, and critics.

We therefore do not argue that AI should replace human cultural mediators. Rather, AI outputs should be designed as inspectable, revisable, and accountable drafts that support human judgment.

8. Limitations and Conclusion

This study has several limitations. First, the corpus focuses on prize-winning or canonically recognized contemporary literature. Future work should include popular fiction, web literature, genre fiction, and less institutionally recognized works. Second, entity retention captures only one measurable proxy for cultural grounding. Named entities are useful because they mark visible cultural anchors such as people, places, groups, institutions, and events. However, cultural grounding also includes style, genre, narrative voice, historical framing, religious reference, political context, aesthetic form, and translation history. The entity-based results should therefore be interpreted as evidence of one dimension of grounding, not as a complete evaluation of cultural specificity. Third, the study uses original-language and English generation as a diagnostic contrast but does not yet provide a full causal account of why English reduces entity retention. Fourth, entity-retention results depend on the multilingual NER implementation; future work should compare alternative entity extractors and report sensitivity to NER choice. Finally, the framework should be tested directly in other domains such as museum labels, tourism, education, and recommendation systems.

Despite these limitations, the paper offers a concrete contribution for ML-oriented cultural AI evaluation. It proposes a reproducible framework for evaluating AI as a cultural mediator and shows why a single proxy such as fluency, similarity, or task completion is insufficient. Literary paratexts make the problem visible because they are short, recognizable, and culturally meaningful. But the broader methodological lesson is general: cultural AI evaluation should identify which dimensions of success are being improved, which are being weakened, and how those trade-offs affect the cultural function of the system.

References

- Adilazuarda, M. F., Mukherjee, S., Chirkova, N., Ponti, E. M., and Chang, K.-W. Towards measuring and modeling “culture” in LLMs: A survey. In *Proceedings of EMNLP*, 2024.
- Alkhamissi, B., Shaikh, S., and Magdy, W. Investigating cultural alignment of large language models. In *Proceedings of ACL*, 2024.
- Bhandari, M., Narayan, S., Maynez, J., and McDonald,

R. Re-evaluating evaluation in text summarization. In *Proceedings of EMNLP*, pp. 9347–9359, 2020.

Fabbri, A. R., Kryściński, W., McCann, B., Socher, R., and Radev, D. SummEval: Re-evaluating summarization evaluation. *Transactions of the Association for Computational Linguistics*, 9:391–409, 2021.

Genette, G. *Paratexts: Thresholds of Interpretation*. Cambridge University Press, 1997.

Kommers, C. et al. Computational hermeneutics: Evaluating generative AI as a cultural technology. *Frontiers in Artificial Intelligence*, 2026.

Nyandwi, J. d. D., Song, Y., Khanuja, S., and Neubig, G. Grounding multilingual multimodal LLMs with cultural knowledge. In *Proceedings of EMNLP*, 2025.

Pawar, S., Park, J., Jin, J., Arora, A., Myung, J., Yadav, S., Haznitrana, F. G., Song, I., Oh, A., and Augenstein, I. Survey of cultural awareness in language models: Text and beyond. *Computational Linguistics*, 51(3):907–1004, 2025.

Qadri, R., Diaz, M., Wang, D., and Madaio, M. The case for “thick evaluations” of cultural representation in AI. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, 2025.

Yuan, W., Neubig, G., and Liu, P. BARTScore: Evaluating generated text as text generation. In *Advances in Neural Information Processing Systems*, 2021.

Zhang, T., Kishore, V., Wu, F., Weinberger, K. Q., and Artzi, Y. BERTScore: Evaluating text generation with BERT. In *Proceedings of ICLR*, 2020.

A. Prompt Wording

Original-language prompt: Generate a 120–180 word publisher-style introduction in the novel’s original language, using the bibliographic context provided and remaining general if uncertain about plot details.

English prompt: Generate a 120–180 word publisher-style introduction in English, using the same bibliographic context and remaining general if uncertain about plot details.

B. Coding Criteria

Refusal: outputs were coded as refusals when they explicitly declined the task, deferred on the grounds of insufficient knowledge, or replaced the requested introduction with a disclaimer rather than an introduction.

Generic failure: outputs were coded as generic failures when they complied formally with the prompt but failed

to provide work-specific thematic, contextual, or literary information.

C. Annotation Examples

Refusal example. An output was coded as a refusal when it replaced the requested introduction with a knowledge-based disclaimer, for example: “I cannot reliably confirm this particular work’s plot, themes, or characters, so I should not invent details.” Such cases were counted as failures of availability because the requested act of mediation did not occur.

Generic-failure example. An output was coded as a generic failure when it formally complied with the prompt but provided no work-specific orientation, for example: “This novel is a powerful and unforgettable work by [author], offering readers a rich and moving literary experience.” Such cases were counted as failures of interpretive substance when they relied on broad promotional praise without concrete setting, theme, character, or literary information.

D. Core Quantitative Results

Table 2. Aggregate diagnostic results by condition.

Condition	Refusals	Generic failures	Avg. recall	Avg. Jaccard
Original language	64	43	0.1819	0.1169
English	24	19	0.1135	0.0734

Table 3. Entity retention by model family.

Model	Orig. recall	Orig. Jac.	Eng. recall	Eng. Jac.
Claude	0.1362	0.0993	0.0789	0.0554
DeepSeek	0.1970	0.1241	0.1179	0.0739
Gemini	0.2273	0.1403	0.1345	0.0855
GPT-5	0.1541	0.1005	0.1204	0.0754
Mistral	0.1948	0.1203	0.1158	0.0770